

# CS342: Operating Systems Lab

Department of Computer Science and Engineering,

Indian Institute of Technology, Guwahati

North Guwahati, Assam 781 039

Training Lesson

**OS Lessons:** Threads, Thread lifecycle, Scheduling, Run-time stack, User Programs (processes)

**Rating:** Easy

**Last update:** 04 June 2018

This document teaches the basic topics to the students of CS342 that are needed by them to begin working on CS342 exercises. These topics will be covered in detail in their CS341 and CS346 classes later. The coverage here is very simple and only enough to understand the exercises.

## Introduction

An operating system is a program just like any C program! It has function `main()`. The translated program gets loaded into a computer memory and runs just like a factorial program may run. When we give command `factorial 10` to a Linux shell, the program code is inserted in the computer memory and control passed to its function `main()`. Each function in a program expects its parameters to be available in the topmost frame (activation record) in an agreed run-time stack.

Function `main()` expects no different. For the command to run factorial program discussed above, function `main()` expects the stack frame to have two parameters. One telling the count of the arguments (`int argc`) and other pointing to an array of pointers (`char *argv[]`) to these argument strings.

However, unlike program `factorial` that will be loaded and supported by a sophisticated operating system like Linux, Linux itself may be loaded and given control by a primitive bootstrap code. `Pintos` is a small but sophisticated operating system. It too gets loaded on its computer; albeit a simulated computer to run. It too gets its initial control through its function `main()`. You can locate this function in file `threads/init.c`

However, `Pintos` (and Linux) are programs that must aid loading and running of the other programs. The operating systems must aid these programs to run and also for reasons of efficient resource utilization let many active programs run concurrently.

Each activity chain defined by execution of a relevant code in `Pintos` is a thread. A relevant piece of code may a program or a utility code within `Pintos` kernel.

All important information for a thread is maintained by `Pintos` in `struct thread` that can be viewed in file `threads/thread.h`. Since many activities are running concurrently

38 there are many concurrent threads in an operating system. However, there is only one  
39 processor. Only one thread can run (progress) at any given point in time.

40 This brings us to a scheduler. Scheduler is a collection of data-structures and functions  
41 (algorithms) within an operating system. Data-structures record all threads in the system.  
42 From time to time, scheduling events occur that cause the scheduler code to be run. When the  
43 code runs it can select a thread from its data-structures to run next. Thread that was running is  
44 inserted into the data-structures to run again at a later time. You will be exploring Pintos  
45 code to find these data-structures and functions as you do your exercises under project  
46 Threads. And add new features and functionality to them.

47 As many threads are running concurrently, one can expect them to be using system resources,  
48 for example, in memory data-structures and external files and devices at the same time.  
49 Inconsistent use of these can cause problems. Thus, you will also find in Pintos code  
50 synchronization primitives that prevent multiple threads from entering critical phases  
51 simultaneously. Only one thread must perform a task at a time that is prone to problems.

## 52 Thread Life-Cycle

53 A thread (activity) may be divided into 5 easily identifiable states. Two of them are obvious –  
54 being created and being removed. To create a thread, the OS kernel must get memory space  
55 for `struct thread` to be the prime representative for the thread. This `struct` is  
56 initialized and then added into the scheduler's data-structures. Once there, it will get access to  
57 the processor time and can do useful activities. The initial activities include location of the  
58 code to run and creation of initial stack-frames.

59 There are three stacks located in `struct thread`! See function `thread_create()`.  
60 Unsurprisingly, these have small space allocations and are for special purposes. The stack  
61 used will be determined by the nature of the activity (functions comprising the activity).  
62 These three stack frames play crucial role in creating a new thread and aiding it to join the  
63 pool of other threads in the system.

64 We will come back to discuss about these stacks embedded in `struct thread` after we  
65 have learned about the stacks and frames (also called activations records for functions) in the  
66 context of a C program.

67 Returning to the issue of thread life cycle, we have already talked of initialization phase and  
68 termination phase. In between the thread is performing useful computation and other related  
69 activities. These activities only occur when the thread is RUN state. It has the processor to  
70 run the instructions in the program code. However, other concurrent threads cannot be  
71 ignored for too long. They too must be run. To do this each RUN state comes with a time  
72 quantum. On completion of that time the thread is put into READY state and some other  
73 thread in ready state is chosen to run. This is done by timer interrupts. At any given time one  
74 would likely see many ready threads and one running threads in the system.

75 However, it may so happen that there is no thread ready to run. Proper functioning of an OS  
76 requires some thread to be running otherwise activities of the system will come to a halt; no  
77 running thread means no system activity! This is easily taken care of by creating a special

78 thread that is always there and ready. It is called *idle* thread. It has the least priority; it runs  
79 only when no one else is around to keep the system going.

80 There is one more important state of the threads. It is state BLOCKED; also called  
81 WAITING. When a thread (program) needs an external event to occur before it can continue  
82 further, for example data read, it must wait for that event. Threads in this state are waiting.  
83 They need not be scheduled by the scheduler to run. When the waited event occurs, the thread  
84 will be transferred back to state ready.

85 Quick summary: All computational and OS activities occur in the context of a thread. The  
86 threads share processor time as disjoint periods. Scheduler ensures that processor time is  
87 shared properly by the threads. Switching (replacing) of a running thread by a ready thread  
88 can be synchronous (usually synchronous switch are voluntary) activity, or asynchronous  
89 (usually involuntary) change. A voluntary switch occurs at the fixed points in the program  
90 when thread/program performs an action that needs to wait for an external (outside the  
91 program) event. In this case, thread normally goes into a block state.

92 An involuntary switch is primarily due to an external or timer interrupt. The thread typically  
93 is placed in a ready queue. The thread will wait for the scheduler to "dispatch" it at a later  
94 time to resume computation.

95 Actions and activities related to thread scheduling and switching are all responsibilities of the  
96 OS. In *Pintos* this code is primarily located in directories *threads/* and *devices/* in  
97 *Pintos*. Application programmers do not directly write any code to support these activities.  
98 Code that an application programmer writes is all related to the application requirements.

## 99 Introduction to Compiler related issues

100 When a function is called there is a set protocol so that the calling function and the called  
101 function can communicate parameters and return values. For this they share a program-wide  
102 stack called runtime stack. Caller creates a new activation record for the function to be called.  
103 In this record or frame the caller places at set and agreed locations space for return value,  
104 function parameters, and the instruction reference at which the caller will resume its activities  
105 after the call completes and returns to the caller.

106 The called function, uses the space above (be warned that the actual direction of stack  
107 growth, on your computer, may be from a low address towards a high address, or it may be  
108 from a high address towards a low address) for its local variables. The called function may, in  
109 turn, call other functions following the exact same protocol.

110 Another issue of interest to us is linking of the functions in the program into a single  
111 executable. We will ignore the dynamic linking and pretend that a monolithic executable is  
112 constructed before the program starts running. This monolith has machine code for all  
113 functions included in the program. Thus, in the executable each calling function knows where  
114 the code for the called function is located. This resolution is done by a linker that is run after  
115 compiler has translated C function codes into machine (or some similar low-level) codes.

116 C also lets us access this location of the function codes as a pointer. You can declare  
117 variables of the right type to hold these pointers. Function name is also such a pointer. Please

118 see function `run_actions()` in file `threads/init.c` to see how the idea is used in  
119 PintOS. Some more information is easily seen on internet; for example:  
120 [https://en.wikipedia.org/wiki/Function\\_pointer#Simple\\_function\\_pointers](https://en.wikipedia.org/wiki/Function_pointer#Simple_function_pointers)

## 121 Test Cases for Exercises in Project Threads

122 The section tries to explain how command `make check` included in PintOS runs for  
123 exercises under project Threads.

124 For projects after thread command `make check` tests PintOS code augmented by you by  
125 creating an executable of a test program and runs the test program on your implementation of  
126 PintOS. The test program is not part of the kernel but a separate program that runs on the  
127 kernel.

128 However, the test programs used in project Threads are included into the kernel by the linker  
129 that runs before PintOS is created. The following command is copied from the screen output  
130 of command `make` under directory `threads/`. Notice that linker `ld` is linking all test  
131 cases here into PintOS kernel

```
132 ld -melf_i386 -T threads/kernel.lds.s -o kernel.o  
133 threads/init.o threads/thread.o threads/switch.o  
134 threads/interrupt.o threads/intr-stubs.o threads/synch.o  
135 threads/palloc.o threads/malloc.o threads/start.o  
136 devices/timer.o devices/kbd.o devices/vga.o devices/serial.o  
137 devices/disk.o devices/input.o  
138 devices/intq.o devices/rtc.o lib/debug.o lib/random.o  
139 lib/stdio.o lib/stdlib.o lib/string.o lib/arithmetic.o  
140 lib/ustar.o lib/kernel/debug.o lib/kernel/list.o  
141 lib/kernel/bitmap.o lib/kernel/hash.o lib/kernel/console.o  
142 tests/threads/tests.o tests/threads/alarm-wait.o  
143 tests/threads/alarm-simultaneous.o  
144 tests/threads/alarm-priority.o tests/threads/alarm-zero.o  
145 tests/threads/alarm-negative.o tests/threads/priority-change.o  
146 tests/threads/priority-donate-one.o  
147 tests/threads/priority-donate-multiple.o  
148 tests/threads/priority-donate-multiple2.o  
149 tests/threads/priority-donate-nest.o  
150 tests/threads/priority-donate-sema.o  
151 tests/threads/priority-donate-lower.o  
152 tests/threads/priority-fifo.o tests/threads/priority-preempt.o  
153 tests/threads/priority-sema.o tests/threads/priority-condvar.o  
154 tests/threads/priority-donate-chain.o  
155 tests/threads/mlfqs-load-1.o tests/threads/mlfqs-load-60.o  
156 tests/threads/mlfqs-load-avg.o tests/threads/mlfqs-recent-1.o  
157 tests/threads/mlfqs-fair.o tests/threads/mlfqs-block.o
```

158  
159 Thus, all test cases in threads project are set as activities within the kernel. These threads are  
160 appropriately called kernel threads – these thread only run the kernel code and runs it in the  
161 supervisory (kernel) mode. The specific (test case) function is identified through an assembly  
162 language code `run_test()`. See file `tests/threads/tests.h`

163 The user programs (all tests cases after exercises on threads), however, must run outside the  
164 kernel in less privileged user mode. These codes are user programs that an unknown user  
165 writes and runs on PintOS kernel. So these are not linked into the kernel code. Instead, they  
166 are loaded into memory by PintOS and allowed to run. The difference in these manners of  
167 code handling is noticed in function `run_task()` in file `threads/init.c`

168 The user programs may seek services from the kernel through system calls – syntactically the  
169 system calls look similar to the function calls but their implementation details are quite  
170 different. The methods and features needed for this are lessons to be learned in the later  
171 projects.

172

173 **Contributors:**

174 Vishv Malhotra, Gautam Barua, Arnab Sarkar