

CS342: Operating Systems Lab

Department of Computer Science and Engineering,

Indian Institute of Technology, Guwahati

North Guwahati, Assam 781 039

Exercise UP-02

OS Lessons: Interrupts, System calls, Exit Status, Return Value

Rating: Moderate

Last updated: 20 February 2017

Please do not start on this exercise until you have successfully completed Exercise UP-01. Attempting to progress without a good understanding of the solutions needed in exercise UP01 is likely to be expensive on your time and efforts.

Task for the Exercise:

System calls are user program's requests to the kernel to perform some service. User programs are prevented from performing these activities directly because their actions may affect the other processes (user programs and even the kernel) running on the system. The kernel provides the coordinating authority to oversee these activities.

A consequence of this restraint is that the kernel must work in a robust isolation from the user program that has requested a service. Thus, kernel and user stacks are separate entities.

At the same time, as a kernel provides the requested services to the user programs it needs to receive the system-call requests and their arguments. Further, the kernel needs a way to communicate the results and service status information back to the user programs (or processes). So we need communication channels between the kernel and the user virtual address spaces. To understand the issues at hand, these sections in PintDoc are definitely your essential readings before you start work on this exercise:

- Section 3.1.5 *Accessing User Memory*,
- Section 3.4.2 *System Calls FAQ*,
- Section 3.5.2 *System Call Details*,
- Section 3.3.2 *Process Termination Messages*,
- Appendix A.4 *Interrupt Handling*,
- Appendix A.4.1 *Interrupt Infrastructure*,
- Appendix A4.2 *Internal Interrupt Handling*

It is important that you take advantage of the program pattern that PintOS code (See function `run_actions()` in file `threads/init.c`) uses to call functions. The pattern maintains an array of pointers to functions. The system call number is used as index in the array to jump to the appropriate function implementing the system call. This is not just efficient, it is less vulnerable to errors and eases the task of adding new functions later. The pattern also keeps the implementation of each system call in a separate function. Function `syscall_handler()` in file `userprog/syscall.c` referred to in Section 3.3.4 *System Calls* becomes a convenient arrival and departure point for all system calls.

Your first and very easy task is to create the skeleton functions for each of the system call mentioned in User Program project and provide a way to call the right function for each system call number for which a request arrives at function `syscall_handler()` in file `userprog/syscall.c`. Each of these functions can initially be set to mirror the older behavior of function `syscall_handler()`.

Goal for the current exercise is to implement and have (at least) the following system calls (see section 3.3.4 *System Calls*) running: `halt()`, `exit()`, and `write()` on `stdout` (`fd = 1`). This will give you at least 10 working test cases from `make check` command.

The `write` system call requires your implementation to return a value back to the user program. The task needs an arrangement that would need a bit of your attention.

System call `halt()` is easy to implement once you work out a small correction you need to make in the specification given on page 29.

System call `exit()` requires whole of the advice given in Section 3.2 *Suggested Order of Implementation* and a bit more.

Our advice is to implement function `process_wait()` in file `process.c` to work as follows:

While the thread associated with `child_tid` is not in state `THREAD_DYING` call `thread_yield()` to avoid spin wait. However, you need a small but subtle change in the provided kernel code to have this part working. This change relates to an allocated resource and when it is safe to release the resource. For the present, we may decide to keep it instead of deallocating the resource. This decision leads to what is described as resource leak as the resource is not being returned to the pool of available resources.

In the last exercise of project *User Program* (UP04) we will revisit the issue of the resource deallocation more earnestly. Unallocated (or leaked) resources degrade the Kernel's ability to host multiple threads/processes over time. Every allocated resource need to be deallocated if the kernel must sustain its vitality indefinitely.

Many a times, PintDoc makes an important suggestion without stressing its importance. The second paragraph after description of system call `close` on page 32 is a good advice. It says that the addresses provided to the system calls by the user program must be carefully checked. PintDoc recommendation is to do these sanitization checks now.

A sanitation check has two parts. First, determine that the address is a valid user virtual address. Second part is to ensure that the address does translate into a real physical address. If either condition is violated, the address is a likely source of trouble for the kernel integrity and stability.

You must create handy functions in file `syscall.c` to test the validity of the addresses passed as arguments in the system calls. The arguments to be tested for valid address and nature of the test depends on the system call. It suffices at this stage to only exercise this test for the system calls you have implemented in this exercise.

You may continue to implement other system calls related to files if you wish. Or you may wait for the next exercise. We recommend that you leave system calls `exec()`, `wait()` for the last exercise. This also may apply to the task described in Section 3.3.5 *Denying Writes to Executables*.

79 Our experiences with `make check`

80 Once again, we have tried to revert our implementation to the stage at the completion of this exercise
81 after we completed the project. Therefore, the tests that are passed and not passed by your
82 implementation of the exercise may not exactly match the results printed below.

```
83 [vmm@progsrv build]$ make check
84 pass tests/userprog/args-none
85 pass tests/userprog/args-single
86 pass tests/userprog/args-multiple
87 pass tests/userprog/args-many
88 pass tests/userprog/args-dbl-space
89 pass tests/userprog/sc-bad-sp
90 pass tests/userprog/sc-bad-arg
91 pass tests/userprog/sc-boundary
92 pass tests/userprog/sc-boundary-2
93 pass tests/userprog/halt
94 pass tests/userprog/exit
95 FAIL tests/userprog/create-normal
96 pass tests/userprog/create-empty
97 FAIL tests/userprog/create-null
98 FAIL tests/userprog/create-bad-ptr
99 pass tests/userprog/create-long
100 FAIL tests/userprog/create-exists
101 FAIL tests/userprog/create-bound
102 FAIL tests/userprog/open-normal
103 FAIL tests/userprog/open-missing
104 FAIL tests/userprog/open-boundary
105 FAIL tests/userprog/open-empty
106 pass tests/userprog/open-null
107 FAIL tests/userprog/open-bad-ptr
108 FAIL tests/userprog/open-twice
109 FAIL tests/userprog/close-normal
110 FAIL tests/userprog/close-twice
111 pass tests/userprog/close-stdin
112 pass tests/userprog/close-stdout
113 pass tests/userprog/close-bad-fd
114 FAIL tests/userprog/read-normal
115 FAIL tests/userprog/read-bad-ptr
116 FAIL tests/userprog/read-boundary
117 FAIL tests/userprog/read-zero
118 pass tests/userprog/read-stdout
119 pass tests/userprog/read-bad-fd
120 FAIL tests/userprog/write-normal
121 FAIL tests/userprog/write-bad-ptr
122 FAIL tests/userprog/write-boundary
123 FAIL tests/userprog/write-zero
124 pass tests/userprog/write-stdin
125 pass tests/userprog/write-bad-fd
126 FAIL tests/userprog/exec-once
127 FAIL tests/userprog/exec-arg
128 FAIL tests/userprog/exec-multiple
129 FAIL tests/userprog/exec-missing
130 pass tests/userprog/exec-bad-ptr
131 FAIL tests/userprog/wait-simple
```

```
132 FAIL tests/userprog/wait-twice
133 FAIL tests/userprog/wait-killed
134 pass tests/userprog/wait-bad-pid
135 FAIL tests/userprog/multi-recurse
136 FAIL tests/userprog/multi-child-fd
137 FAIL tests/userprog/rox-simple
138 FAIL tests/userprog/rox-child
139 FAIL tests/userprog/rox-multichild
140 pass tests/userprog/bad-read
141 pass tests/userprog/bad-write
142 pass tests/userprog/bad-read2
143 pass tests/userprog/bad-write2
144 pass tests/userprog/bad-jump
145 pass tests/userprog/bad-jump2
146 FAIL tests/userprog/no-vm/multi-oom
147 FAIL tests/filesys/base/lg-create
148 FAIL tests/filesys/base/lg-full
149 FAIL tests/filesys/base/lg-random
150 FAIL tests/filesys/base/lg-seq-block
151 FAIL tests/filesys/base/lg-seq-random
152 FAIL tests/filesys/base/sm-create
153 FAIL tests/filesys/base/sm-full
154 FAIL tests/filesys/base/sm-random
155 FAIL tests/filesys/base/sm-seq-block
156 FAIL tests/filesys/base/sm-seq-random
157 FAIL tests/filesys/base/syn-read
158 FAIL tests/filesys/base/syn-remove
159 FAIL tests/filesys/base/syn-write
160 47 of 76 tests failed.
161 make: *** [check] Error 1
162 [vmm@progsrv build]$
163
```

164 **Contributing Authors:**

165 **Vishv Malhotra, Gautam Barua, Rashmi Dutta Baruah**